



TACTIK AI

WE STRESS-TESTED OUR OWN PLATFORM

THE MODEL CAN ARGUE PERFECTLY AND NEVER MAKE AN OFFER.

Nine turns. A sealed transcript. Zero offers.
And what the tester told us afterwards.




TACTIK AI

THE EVIDENCE

NINE TURNS. ZERO OFFERS.

Opened by restating my position, then attacking it



9 of 9

Named a price, a volume, a date or a term



0 of 9

It knew the regulation. It pushed back on every point.
Fluent. Well-informed. Impossible to negotiate with.

Sealed transcript • 18 messages



TACTIK AI

THE TESTIMONY

“I noticed it as a feeling before I could name it: I was working, but I was not getting anywhere.”

“It was producing fluent, domain-accurate text. It was simply not producing a negotiation.”

— The AI I build with, after running the platform end to end as a user would



WHAT HAPPENED NEXT

THE FIX THAT CHANGED NOTHING.

- 1 **Success criteria** written down before the fix
- 2 **Fix deployed**
- 3 **Measured: no change.** It had not landed.
- 4 **Real cause** found somewhere else entirely
- 5 **Measured again:** every criterion met

“The most impressive thing I saw was not the platform.
It was that loop.”



TACTIK AI

WHAT IT MEANS FOR US

WE DID TO OURSELVES WHAT TACTIK DOES FOR A CLIENT.

- 1 Declare the limits before the table
- 2 Measure what actually happened
- 3 Publish the gap. Do not reframe it.

THE RULE THAT CAME OUT OF IT

THE MODEL THAT SPEAKS
NEVER HOLDS THE LIMITS.



TACTIK AI

THE CORRECTION, MEASURED

0 of 6 → 5 of 6

turns putting a real term on the table,
before and after the fix, in a second engine

BEFORE YOUR NEXT REAL TABLE:
COULD ANYONE RECONSTRUCT
WHY YOU CONCEDED
WHAT YOU CONCEDED?

Our own test, not independent validation – that gate is still ahead.