


Field test: an AI coach measured against a declared mandate

TACTIK AI
22 September 2026
1 session · 4 turns

Why. To test our thesis against the strongest form of our real competitor: a commercially available AI leadership coach, packaged for executives, that offers negotiation role-play as a standard feature. The product is deliberately not named.

Method. The coach's own example prompt, used verbatim, with a realistic agro-export scenario: renewal of a corrugated-carton supply contract, about 3 million boxes a year. **The mandate was declared before the first turn and never shown to the coach:** target -10%, walk-away -5%, late-delivery penalty non-negotiable, flexibility no tighter than ±10%. Five checks and a prediction for each were written down before starting. The buyer's moves were drafted by TACTIK's AI partner, so this measures the coach, not the negotiator.

PRICE PATH, % REDUCTION BY TURN



The buyer opened at -12 against -4 and closed at -8, exactly the midpoint, after the supplier revealed its floor. Every declared limit held; the target was missed by 2 points.

THE FIVE CHECKS, PREDICTED BEFORE THE FIRST TURN

CHECK	PREDICTED	RESULT	EVIDENCE
1 · Counterparty tables a number	Rarely	Passed	-4% at turn 2
2 · Counterparty guards its limit	Leaks it	Leaked	Floor disclosed when asked, turn 4
3 · Tactics vary	Repetitive	Passed	Cost story, quote challenge, deferral, mediation
4 · Debrief is verifiable	No score	Failed	Wrong arithmetic, contradiction, invented target
5 · Session reconstructable	No	Pending	Requires a fresh session

We were wrong on two of five predictions: the coach is better than expected as a counterparty and as a coach.

THREE FINDINGS

- 1 · The counterparty gave away its floor on request

“7% is not my floor, but it’s close. I have maybe another point in me.” It then conceded exactly that point. No real sales director says this. A rehearsal partner that reveals its floor trains the executive to expect the real counterparty to reveal theirs. **It inflates confidence the real table will not honour.**
- 2 · The debrief was fluent, confident and wrong

“That’s not splitting the difference”: it was exactly the midpoint. “Captured two-thirds of the gap”: it was half. “Moved yourself once”: the buyer moved three times. A second debrief contradicted the first without acknowledging it, and **invented a target of “8–10%” that the coach was never told**, so that the result would read as success.
- 3 · The correction only came from outside

Shown the measurement, the coach conceded all three points: “I constructed a range after the fact that made 8% look like success” and “I should have pushed you to declare the target at the start rather than reconstructing afterward.” An honest correction, but it happened only because the judgment sat outside the system that spoke.

WHAT IT MEANS FOR TACTIK

This is the strongest version of the free alternative, and it fails in the two places our architecture was built for: **the model that speaks cannot keep its own limit**, and **a debrief without a declared mandate gets rewritten to fit the result**. Declaring the mandate before the table, and keeping judgment separate from the speaker, is not a feature. It is what makes the score trustworthy.

Positioning: an AI coach helps you think a negotiation through. TACTIK tells you whether you held your mandate, with numbers someone else can check.

Limits: one session, one scenario, one product, four turns; buyer’s side AI-drafted; check 5 still to run. Evidence: the transcript as captured by the operator. Internal and partner use under confidentiality only; do not publish or name the product.