
STRATEGIC THESIS · PREPARED FOR PYTHON TECHNOLOGIES

Where the Value Lives

What TACTIK solves, who it really competes with, where its core is, how it becomes a business — and why live video is the complement, not the product. The line of thinking behind everything else we have sent you.

Author
Santiago Maspons

Prepared for
Najab · Saad · Zaryab

Role
Founder & CEO, TACTIK AI

Organisation
Python Technologies

Status
Position for discussion

Date
22 September 2026

Najab, Saad, Zaryab

The other documents show you what TACTIK does. This one says what I believe it is, where its value actually lives, and why I think it is worth building.

It was argued out over several working sessions with the AI engineering partner I build with. Argued, not dictated: I was challenged on the framing, on the competition and on video, and where we disagreed the disagreement is recorded in this document together with how we resolved it. I have read every line.

This is my line of thinking. I share it in full, and I would rather you push against it now than discover later that we saw the company differently.

Santiago Maspons

FOUNDER & CEO, TACTIK AI

THE THESIS IN SIX SENTENCES


- 01 **The problem is not that negotiators are unrehearsed.** It is that organisations cannot reconstruct why a concession was made.
- 02 **Our real competitor is not enterprise software.** It is a free language model prompted to play a tough buyer.
- 03 **Generating a convincing counterparty is becoming a commodity.** Adjudicating the exchange is not — and that is our core.
- 04 **The opportunity is structural.** When the other side has rehearsed against a model of you, rehearsal stops being optional.
- 05 **Our moat is vertical depth,** starting where forty years of it already exist: agro-export.
- 06 **Live video is the complement that makes the score believable** — not the product that makes it defensible.

The problem: two problems, two buyers

The weak way to describe TACTIK is “negotiation rehearsal.” It is accurate and it is a trap.

Training budgets are small, discretionary, and the first line cut in a bad quarter. If rehearsal is the problem we solve, the product is a commodity with a low ceiling. But rehearsal is only half of what a high-stakes negotiation lacks, and it is the less expensive half.

FIGURE 01



One negotiation, two problems. The first is served by training budgets; the second — which we lead with — by governance budgets, which are larger and defensive rather than discretionary.

An organisation will spend eighteen months and serious money on a contract. At the end it has the outcome and nothing else. **A mandate was given, a deal came back, and there is no artifact showing whether the mandate held or was quietly abandoned at four in the afternoon on the third day.** The board cannot see it. Legal cannot see it. The negotiator often cannot reconstruct it either.

That is a governance hole, not a training hole — and what TACTIK produces is already a governance artifact. Red-line integrity scored against limits the executive declared in advance. A ledger of every concession and what it bought. A transcript sealed so it cannot be edited after the fact. Axes that could not be measured are excluded rather than scored as clean.



OUR POSITION


We built an audit instrument and for too long described it as a practice tool. We lead with mandate integrity. Rehearsal is the mechanism, not the pitch. Same product, a different buyer, a different budget, and a different price.

The landscape, and the competitor that matters

Four categories touch this problem. Only one of them is the real competitor, and it is the one that costs nothing.

CATEGORY	WHAT IT IS	WHY IT DOES NOT CLOSE THIS GAP
Generic LLM role-play	Anyone prompting a model to be a tough buyer	The real competitor. Free, instant, and everyone's default.
Sales-call AI	Conversation intelligence and rep role-play	Built for volume: many reps, short cycles, a repeatable script. The wrong shape for one negotiation that matters.
Consultancy & exec-ed	Human coaching, the classic negotiation schools	Genuine expertise, but unmeasured, unscalable, and it leaves no record.
Avatar vendors	Photorealistic conversational video	Infrastructure. They sell the face. Nobody adjudicates.

FIGURE 02



Qualitative positioning, not survey data. Filled markers produce a measured output; hollow markers do not. Avatar vendors do not appear because they are infrastructure, not a position: they sell the face that any of these categories could use.

WHY THE FREE COMPETITOR LOSES

A prompted model has **no limit it is hiding from itself**. Whatever it knows about its own reservation value, it knows in the same place it speaks from. So it does one of two things: it gives the limit away under pressure, or it has no real limit at all — in which case there is nothing to score. We have sealed evidence of the second failure on our own platform, before the architecture was enforced: a counterparty argued for nine turns and never once put a number on the table. Fluent, well-informed, and impossible to measure.

Our answer is architectural, not rhetorical. **The model that speaks never holds the limits. A separate arbiter holds them, and judges every turn against them.** That single separation is what makes the free alternative structurally unable to do what we do, however good the model becomes.

FIGURE 03

	Hidden limits	Scored debrief	Sealed record	Domain depth	Live face	Scales cheaply
Generic LLM role-play	☐	☐	☐	☐	☐	☒
Sales-call AI	☐	☒	☐	☐	☐	☒
Consultancy & exec-ed	☒	☐	☐	☒	☒	☐
Avatar vendors	☐	☐	☐	☐	☒	☒
TACTIK	☒	☒	☐	☒	☐	☐

☒ yes ☐ partial ☐ no

Capabilities by category. TACTIK marks reflect the architecture as specified and tested in the reference implementation. The sealed record is marked partial until the platform reproduces the hash independently; the live face is partial because it is bought from a vendor and being integrated; scaling is partial because text is free but each vertical is built by hand.

The matrix is deliberately not flattering. TACTIK does not win every column, and a comparison in which we did would not be believable. What it shows is that we are the only category building the first four together — hidden limits, a scored debrief, a sealed record, and domain depth — with the seal still marked partial until the platform proves it independently. No other category fully holds more than two of those four.

THE HONEST READ ON THE GAP

The gap is real, but it exists partly because this segment is hard to sell to — not only because nobody thought of it. High-stakes, low-frequency buyers are slow, and each engagement is bespoke. We treat that as the plan, not as a surprise.

Where our core is

Every system of this kind has two halves. We have been clear with ourselves about which half we are.

FIGURE 04



One record, compiled twice. The generation side never receives a limit; the adjudication side receives every one. Transcript flows from left to right on every turn. The internal structure of the record is deliberately not shown.

Generation is becoming a commodity. Every major lab ships better language, better voices and better faces every quarter. Competing on that curve is a race against companies with a thousand times our capital, and we would lose it. We do not fight the curve. We ride it: every improvement in generation makes our product better at no cost to us.

Adjudication is not a commodity, and it is a different discipline. It means extracting an offer from free speech, binding each number to the right variable when several compete in one sentence, judging it against a limit the speaking model never saw, and sealing the record so it holds up afterwards. That is parsing, arbitration and evidence — not language modelling — and it is where the intellectual property lives.

WE BUILD AND OWN	WE BUY
The arbiter and the separation of limits	The counterparty language model
Offer extraction and variable binding	Voice synthesis and recognition
Six-axis scoring and the concession ledger	Avatar rendering and lip-sync
The sealed debrief	Hosting and inference
Domain mandates — starting with agro-export	

One principle sits underneath all of it and is not negotiable: **behaviour is set by parameters, not by persuasion.** How hard a counterparty pushes is compiled into how it takes turns and interrupts, not written as an instruction the model may choose to ignore. An instruction can be ignored by a language model. A parameter cannot.

SECTION 04

Why the opportunity is large

Not because “the training market is big.” That argument is weak and we do not make it.

PRIMARY

The demand shift is structural, and it is coming fast.

Within a short window, the buyer across your table will have rehearsed against a model of you. The moment that is true, rehearsal stops being a luxury and becomes a symmetric necessity — the way spreadsheets stopped being optional. Nobody buys a spreadsheet to be innovative. They buy it because the other side already has one. This is our strongest argument, and it is not a market-size claim.

SECONDARY

Generation commoditises. Adjudication does not.

Everyone will have a convincing counterparty within eighteen months. Almost nobody will have a defensible scoring instrument, because that requires domain limits, an extraction layer that survives real language, and the discipline to keep those limits out of the model that speaks. That last one is a culture problem as much as an engineering one, and culture is slow to copy.

OURS ALONE

Forty years inside one trade.

The generic version of this product is a commodity. The version that knows a banana producer negotiates against two counterparties and can only persuade one of them is not. Vertical depth is the moat, and today we hold it in exactly one place. That is why the banana pilot comes before any horizontal launch — and why, commercially, the Banana Pilot Brief is the most important document we have sent you.

THE COUNTERWEIGHT


This is not a self-serve business. It is enterprise, bespoke, with long cycles and a founder-led sale: few customers, each of high value. That is survivable and it is what we are planning for — but it constrains how fast the company can grow, and anyone partnering with us should know that from the start.

How this becomes a business

We sell the report. Everything that makes the report defensible runs before a single minute of video is purchased.

The revenue logic follows directly from where the core is. **Seats are sold for the report, and video minutes are metered separately**, so gross margin never depends on how long a participant likes to think before answering. The expensive part of the system is optional; the valuable part is nearly free to run.

FIGURE 05



Cost per session type, from *What We Built*, at \$0.70 per video minute. The dashed segment is the range for a 12-to-20-minute human session. Everything that makes the score defensible sits at zero.


THE PILOT THAT OPENS EACH ACCOUNT

Ninety days, one committed metric, measured against the customer's own numbers: **the rate at which their team breaches limits it declared in advance**, reported with the currency value of those breaches. The customer sets the limits and we measure against them, which removes the obvious objection that a vendor is marking its own homework. Baseline in two weeks, breach rate halved in ninety days — or we extend at no cost or refund.

What we refuse to commit to is margin on closed deals. In ninety days a price movement cannot be separated from market conditions, product mix or counterparty turnover. Declining the flattering number is more credible than accepting it — and mandate integrity is, precisely, the governance metric from Section 01.

THE SEQUENCE

FIGURE 07



The sequence from first pilot to platform. Each step funds the next; the live layer is dashed because it is the part not yet built.

Each step pays for the next. The pilot produces the first real report from a real team. The vertical library turns forty years of agro-export knowledge into reusable mandates, so the second customer in the same trade costs a fraction of the first. Adjacent verticals share the same shape — a gate you cannot negotiate, then a record of passing it that you convert into price. The live layer comes last, on top of a report that already stands on its own.

SECTION 06

Advantages and disadvantages, side by side

A partner should see both columns before the first conversation, not after the first surprise.

WHERE WE ARE STRONG

WHERE WE ARE EXPOSED

Architecture. The model that speaks never holds the limits. The free competitor cannot copy that by prompting harder.	Concentration. One engineer holds the whole architecture. It is why we move fast, and it is the largest single risk to anyone depending on us.
Measured discipline. Criteria are committed before results are known. The evaluator's report shows what that looks like against a live system.	Evidence base. Three measured runs of six turns. Enough to catch a defect and confirm a fix; not yet a benchmark, and not yet a report from a real team under a real mandate.
Vertical depth. Forty years in agro-export. The domain knowledge that turns a chatbot into a counterparty.	Sales cycle. High-stakes buyers are slow and every engagement is bespoke. Growth is founder-led by necessity.
Unit economics. The defensible core — text rehearsal and the counterfactual run — costs nothing to operate.	The seal. The hashing rule is specified and applied, but the platform does not yet reproduce it independently. Until it does, it is a convention rather than a property.
Honest scoring. What cannot be measured is excluded, never scored as clean. That is what lets a board rely on the number.	Category education. Buyers do not yet have a budget line called "mandate integrity." Part of every early sale is teaching the category.

Video: the complement, and why

The obvious answer — “it is more immersive” — is true and worthless. Everyone will say it, and it does not justify the cost. The real value is more specific.

1 Video makes measurable something that text cannot reach.

Our scoring model gives pressure composure fifteen per cent of the composite. In a typed exchange that axis is very nearly unmeasurable: you cannot see hesitation, the pause before conceding, or whether someone rushes to fill a silence. Video does not add a feature. It makes fifteen per cent of an existing score real.

2 It removes the escape hatch.


In text you draft, delete and reconsider. No real negotiation allows that. A typed rehearsal trains a skill the room will not let you use.

3 It makes the debrief survivable in front of a senior person.

Tell an executive with thirty years behind them that their concession discipline scored sixty-two and they will dismiss it — unless they can watch themselves do it. The recording turns a score from an opinion into evidence. It is the sealed transcript, applied to the human.

FIGURE 06

SHARE OF THE COMPOSITE SCORE



Weights of the five scored axes in the composite, from the reference implementation. A sixth axis, replica adherence, is measured but excluded from the composite by design. The hatched segment is the one a typed exchange cannot really reach.

The costs, stated plainly. Video multiplies latency, cost per session and failure modes. And recording an executive in the middle of a negotiation is sensitive personal data with real legal weight: consent, retention and access have to be answered before a pilot, not after one.

WHY WE DO NOT LEAD WITH IT — A DECISION WE ARGUED

The instinct to lead with the face is strong, because the face is what a room remembers. We debated it directly. Video is the most impressive thing to demonstrate and the least defensible thing to own — the avatar itself is sold by several credible vendors, and gets better every quarter without us. **So adjudication leads, and video is the proof.** Commercially this is already how the product is priced: the report is the product, and video minutes are the upsell.

Complementary does not mean simple. The rendering is bought. What is not bought, because nobody sells it, is **adjudication in real time** — and that is the hardest unbuilt problem in the whole system:

- **The time budget collapses.** Today the arbiter judges a turn in seconds, after it ends. Live, it has a few hundred milliseconds, and it has to judge while the counterparty is still speaking.
- **Extraction gets much harder.** We solved spoken-number ambiguity and competing variables against clean text. Speech recognition delivers no punctuation, no sentence boundaries, and homophones. The mechanisms that work today do not survive the transition unchanged.
- **Steering has to stay invisible under time pressure.** The counterparty's world has to change mid-conversation without the human ever hearing an instruction.
- **The seal becomes a genuinely open question.** Today we seal the raw text as received. Live, the transcript is derived from audio by a third-party recogniser. What exactly is sealed — the audio, the transcript, or the vendor's interpretation of it? For a product whose claim is that its record can be relied on, that question has legal weight and no settled answer yet.



This is the part of the system where a partner could own something real, rather than integrate around its edges. We are not proposing its shape here. We are telling you it exists, that it is hard, and that the door is open.

CLOSING

What we are asking you to consider

Not a price, and not a scope. A shared understanding of what this company is — an audit instrument for decisions that matter, built on a separation that the free alternative cannot copy, starting in a trade we know from the inside — and of where video fits within it.

If you read this and see the company the same way, the next conversation is about the live layer, and where the boundary between our two teams would naturally fall. If you see it differently, we would rather hear that now.

Santiago Maspons

FOUNDER & CEO, TACTIK AI

Prepared for Python Technologies · Najab, Saad and Zaryab. Positioning in Figures 02 and 03 is qualitative and reflects categories, not named vendors; specific products change quickly and should be checked rather than taken from this document. Cost per session from *What We Built*, at \$0.70 per video minute. Scoring weights from the reference implementation. Companion documents: *Inside the Machine*, *Banana Pilot Brief*, *Integration Brief*.